


Machine Learning 1

#Clustering Methods

kmeans clustering in R is done with the 'kmeans()' function. Here we make up some data to test and learn with. #rev command reverses the data set #cbind binds vectors by column. rbind binds them by rows


```
tmp <- c(rnorm(30, 3), rnorm(30, -3))
data <- cbind(x = tmp, y=rev(tmp))
plot(data)
```



Run 'kmeans()' set K to 2 nstrt 20. The thing with kmeans is you have to tell it how many clusters you want

```
km <- kmeans(data, centers = 2, nstart=20)
km
```

```
## K-means clustering with 2 clusters of sizes 30, 30
##
## Cluster means:
```

#hclust - hierarchical clustering We will use the “hclust()” function on the same data as before and see how this method works


```
hc <- hclust( dist(data))
hc
```

```
##
## Call:
## hclust(d = dist(data))
##
## Cluster method   : complete
## Distance         : euclidean
## Number of objects: 60
```

hclust has a plot method

```
plot(hc)
```

Cluster Dendrogram



```
dist(data)
hclust (*, "complete")
```


To find our membership vector we need to “cut” the tree and for this we use the ‘`cutree()`’ function and tell it the height to cut at.

```
cutree(hc, h=7)
```

[illegible]

```
plot(hc)
```

Cluster Dendrogram



```
dist(data)
hclust(*, "complete")
```

```
abline
```


```
## function (a = NULL, b = NULL, h = NULL, v = NULL, reg = NULL,
##   coef = NULL, untf = FALSE, ...)
## {
##   int_abline <- function(a, b, h, v, untf, col = par("col"),
##     lty = par("lty"), lwd = par("lwd"), ...) .External.graphics(C_abline,
##     a, b, h, v, untf, col, lty, lwd, ...)
##   if (!is.null(reg)) {
##     if (!is.null(a))
##       warning("'a' is overridden by 'reg'")
##     a <- reg
##   }
##   if (is.object(a) || is.list(a)) {
##     p <- length(coefa <- as.vector(coef(a)))
##     if (p > 2)
##       warning(gettextf("only using the first two of %d regression coefficients",
##         p), domain = NA)
##     islm <- inherits(a, "lm")
##     noInt <- if (islms)
##       !as.logical(attr(stats::terms(a), "intercept"))
##     else p == 1
##     if (noInt) {
##       a <- 0
##       b <- coefa[1L]
##     }
##   }
```

```
##         else {
##             a <- coefa[1L]
##             b <- if (p >= 2)
##                 coefa[2L]
##             else 0
##         }
##     }
##     if (!is.null(coef)) {
##         if (!is.null(a))
##             warning("'a' and 'b' are overridden by 'coef'")
##         a <- coef[1L]
##         b <- coef[2L]
##     }
##     int_abline(a = a, b = b, h = h, v = v, untf = untf, ...)
##     invisible()
## }
## <bytecode: 0x7fc15b0ba758>
## <environment: namespace:graphics>
```

We can also use 'cutree()' and state the number of k clusters we want...

```
grps <- cutree(hc, k=2)
```

```
plot(data, col=grps)
```



#principal component analysis (PCA)

PCA is a useful analysis method when you have lots of dimensions in your data...

##PCA of UK food data

Import the data from a CSV file

```
url <- "https://tinyurl.com/UK-foods"
x <- read.csv(url)
```

How many row and cols?

```
dim(x)
```

```
## [1] 17 5
```

#remove first col as this is not a col just the rows but it is seeing it as a col

```
x[,1]
```

```
## [1] "Cheese"          "Carcass_meat "      "Other_meat "
## [4] "Fish"             "Fats_and_oils "     "Sugars"
## [7] "Fresh_potatoes "   "Fresh_Veg "         "Other_Veg "
## [10] "Processed_potatoes " "Processed_Veg "      "Fresh_fruit "
## [13] "Cereals "          "Beverages"          "Soft_drinks "
## [16] "Alcoholic_drinks " "Confectionery "
```

```
rownames(x) <- x[,1]
x <- x[,-1]
x
```


```
##           England Wales Scotland N.Ireland
## Cheese           105   103      103        66
## Carcass_meat     245   227      242       267
## Other_meat       685   803      750       586
## Fish            147   160      122        93
## Fats_and_oils    193   235      184       209
## Sugars           156   175      147       139
## Fresh_potatoes   720   874      566      1033
## Fresh_Veg        253   265      171       143
## Other_Veg        488   570      418       355
## Processed_potatoes 198   203      220       187
## Processed_Veg     360   365      337       334
## Fresh_fruit     1102  1137      957       674
## Cereals          1472  1582     1462      1494
## Beverages         57    73        53        47
## Soft_drinks     1374  1256     1572      1506
## Alcoholic_drinks  375   475      458       135
## Confectionery     54    64        62        41
```

#below is a better way to do this. above way - you will lose a col everytime you run it


```
url <- "https://tinyurl.com/UK-foods"
x <- read.csv(url, row.names = 1)
x
```

```
##           England Wales Scotland N.Ireland
## Cheese           105   103      103       66
## Carcass_meat     245   227      242      267
## Other_meat       685   803      750      586
## Fish            147   160      122       93
## Fats_and_oils    193   235      184      209
## Sugars           156   175      147      139
## Fresh_potatoes   720   874      566     1033
## Fresh_Veg        253   265      171      143
## Other_Veg        488   570      418      355
## Processed_potatoes 198   203      220      187
## Processed_Veg     360   365      337      334
## Fresh_fruit     1102  1137      957      674
## Cereals          1472  1582     1462     1494
## Beverages         57    73       53       47
## Soft_drinks     1374  1256     1572     1506
## Alcoholic_drinks  375   475      458      135
## Confectionery     54    64       62       41
```


```
barplot(as.matrix(x), col=rainbow(17))
```




```
barplot(as.matrix(x), col=rainbow(17))
```



```
mycols <- rainbow(nrow(x))  
pairs(x, col=mycols, pch=16)
```



##PCA to the rescue

Here we will use the base R function for PCA, which is called 'prcomp()' First transpose the data with "t" because 'prcomp()' wants the the rows. It says this in the help page

t(x)

```
##      Cheese Carcass_meat Other_meat Fish Fats_and_oils Sugars
## England      105         245      685  147             193  156
## Wales         103         227      803  160             235  175
## Scotland      103         242      750  122             184  147
## N.Ireland       66         267      586   93             209  139
##      Fresh_potatoes Fresh_Veg Other_Veg Processed_potatoes
## England           720        253      488                 198
## Wales             874        265      570                 203
## Scotland          566        171      418                 220
## N.Ireland         1033        143      355                 187
##      Processed_Veg Fresh_fruit Cereals Beverages Soft_drinks
## England           360        1102    1472          57     1374
## Wales             365        1137    1582          73     1256
## Scotland          337         957    1462          53     1572
## N.Ireland          334         674    1494          47     1506
##      Alcoholic_drinks Confectionery
## England              375           54
## Wales                475           64
## Scotland             458           62
## N.Ireland            135           41
```

```
prcomp(x)
```

```
## Standard deviations (1, .., p=4):  
## [1] 919.13914 132.06254 88.58981 28.62021  
##  
## Rotation (n x k) = (4 x 4):  
##           PC1      PC2      PC3      PC4  
## England  0.4901572 -0.2808420 0.07019969 0.82215915  
## Wales    0.4981704 -0.3475970 0.64010493 -0.47039166  
## Scotland 0.5042356 -0.2482197 -0.76259165 -0.32029268  
## N.Ireland 0.5072659 0.8594714 0.06157655 -0.01409337
```

```
pca <- prcomp( t(x))  
summary(pca)
```

```
## Importance of components:  
##           PC1      PC2      PC3      PC4  
## Standard deviation 324.1502 212.7478 73.87622 4.189e-14  
## Proportion of Variance 0.6744 0.2905 0.03503 0.000e+00  
## Cumulative Proportion 0.6744 0.9650 1.00000 1.000e+00
```


We want score plot (a.k.a PCA plot). Basically of PC1 vs PC2

```
attributes(pca)
```

```
## $names  
## [1] "sdev"      "rotation" "center"   "scale"    "x"  
##  
## $class  
## [1] "prcomp"
```


We are after the `pca$x` component for this plot

```
plot(pca$x[,1:2])  
text(pca$x[,1:2], labels = colnames(x))
```



We can also examine the PCA “loadings”, which tell us how much original variables contribute to each new PCA...

```
barplot(pca$rotation[,1], las=2)
```



RNASEQ

```
url2 <- "https://tinyurl.com/expression-CSV"
rna.data <- read.csv(url2, row.names=1)
head(rna.data)
```

```
##      wt1 wt2 wt3 wt4 wt5 ko1 ko2 ko3 ko4 ko5
## gene1 439 458 408 429 420 90 88 86 90 93
## gene2 219 200 204 210 187 427 423 434 433 426
## gene3 1006 989 1030 1017 973 252 237 238 226 210
## gene4 783 792 829 856 760 849 856 835 885 894
## gene5 181 249 204 244 225 277 305 272 270 279
## gene6 460 502 491 491 493 612 594 577 618 638
```

```
nrow(rna.data)
```

```
## [1] 100
```

```
ncol(rna.data)
```

```
## [1] 10
```

```
colnames(rna.data)
```

```
## [1] "wt1" "wt2" "wt3" "wt4" "wt5" "ko1" "ko2" "ko3" "ko4" "ko5"
```

```
pca.rna <- prcomp(t(rna.data), scale=TRUE)
summary(pca.rna)
```

```
## Importance of components:
```

```
##              PC1      PC2      PC3      PC4      PC5      PC6      PC7
## Standard deviation  9.6237 1.5198 1.05787 1.05203 0.88062 0.82545 0.80111
## Proportion of Variance 0.9262 0.0231 0.01119 0.01107 0.00775 0.00681 0.00642
## Cumulative Proportion 0.9262 0.9493 0.96045 0.97152 0.97928 0.98609 0.99251
##              PC8      PC9      PC10
## Standard deviation  0.62065 0.60342 3.348e-15
## Proportion of Variance 0.00385 0.00364 0.000e+00
## Cumulative Proportion 0.99636 1.00000 1.000e+00
```

#summary above shows that PC2 in the plot below is not as significant as PC1. i.e. it is the short line on the graph

```
plot(pca.rna$x[,1:2])
text(pca.rna$x[,1:2], labels = colnames(rna.data))
```

