

Class15_RNASeq

Caitriona Brennan

11-17-21

```
library(BiocManager)
```

```
library(DESeq2)
```

```
## Loading required package: S4Vectors
```

```
## Loading required package: stats4
```

```
## Loading required package: BiocGenerics
```

```
##
```

```
## Attaching package: 'BiocGenerics'
```

```
## The following objects are masked from 'package:stats':
```

```
##
```

```
##      IQR, mad, sd, var, xtabs
```

```
## The following objects are masked from 'package:base':
```

```
##
```

```
##      anyDuplicated, append, as.data.frame, basename, cbind, colnames,  
##      dirname, do.call, duplicated, eval, evalq, Filter, Find, get, grep,  
##      grepl, intersect, is.unsorted, lapply, Map, mapply, match, mget,  
##      order, paste, pmax, pmax.int, pmin, pmin.int, Position, rank,  
##      rbind, Reduce, rownames, sapply, setdiff, sort, table, tapply,  
##      union, unique, unsplit, which.max, which.min
```

```
##
```

```
## Attaching package: 'S4Vectors'
```

```
## The following objects are masked from 'package:base':
```

```
##
```

```
##      expand.grid, I, unname
```

```
## Loading required package: IRanges
```

```
## Loading required package: GenomicRanges
```

```
## Loading required package: GenomeInfoDb
```

```

## Loading required package: SummarizedExperiment

## Loading required package: MatrixGenerics

## Loading required package: matrixStats

##
## Attaching package: 'MatrixGenerics'

## The following objects are masked from 'package:matrixStats':
##
##   colAlls, colAnyNAs, colAnys, colAvgPerRowSet, colCollapse,
##   colCounts, colCummaxs, colCummins, colCumprods, colCumsums,
##   colDiffs, colIQRDiffs, colIQRs, colLogSumExps, colMadDiffs,
##   colMads, colMaxs, colMeans2, colMedians, colMins, colOrderStats,
##   colProds, colQuantiles, colRanges, colRanks, colSdDiffs, colSds,
##   colSums2, colTabulates, colVarDiffs, colVars, colWeightedMads,
##   colWeightedMeans, colWeightedMedians, colWeightedSds,
##   colWeightedVars, rowAlls, rowAnyNAs, rowAnys, rowAvgPerColSet,
##   rowCollapse, rowCounts, rowCummaxs, rowCummins, rowCumprods,
##   rowCumsums, rowDiffs, rowIQRDiffs, rowIQRs, rowLogSumExps,
##   rowMadDiffs, rowMads, rowMaxs, rowMeans2, rowMedians, rowMins,
##   rowOrderStats, rowProds, rowQuantiles, rowRanges, rowRanks,
##   rowSdDiffs, rowSds, rowSums2, rowTabulates, rowVarDiffs, rowVars,
##   rowWeightedMads, rowWeightedMeans, rowWeightedMedians,
##   rowWeightedSds, rowWeightedVars

## Loading required package: Biobase

## Welcome to Bioconductor
##
##   Vignettes contain introductory material; view with
##   'browseVignettes()'. To cite Bioconductor, see
##   'citation("Biobase")', and for packages 'citation("pkgname)".

##
## Attaching package: 'Biobase'

## The following object is masked from 'package:MatrixGenerics':
##
##   rowMedians

## The following objects are masked from 'package:matrixStats':
##
##   anyMissing, rowMedians

```

Today we examine a published RNA-seq experiment where airway smooth muscle cells were treated with dexam

We need tow things : -1:count data 2: col data

```
counts <- read.csv("airway_scaledcounts.csv", row.names = 1)
metadata <- read.csv("airway_metadata.csv")
```

```
head(counts)
```

```
##                SRR1039508 SRR1039509 SRR1039512 SRR1039513 SRR1039516
## ENSG00000000003          723          486          904          445          1170
## ENSG00000000005           0           0           0           0           0
## ENSG000000000419         467          523          616          371          582
## ENSG000000000457         347          258          364          237          318
## ENSG000000000460          96           81           73           66          118
## ENSG000000000938           0           0           1           0           2
##                SRR1039517 SRR1039520 SRR1039521
## ENSG00000000003         1097          806          604
## ENSG00000000005           0           0           0
## ENSG000000000419         781          417          509
## ENSG000000000457         447          330          324
## ENSG000000000460          94          102           74
## ENSG000000000938           0           0           0
```

```
head(metadata)
```

```
##           id      dex celltype    geo_id
## 1 SRR1039508 control   N61311 GSM1275862
## 2 SRR1039509 treated   N61311 GSM1275863
## 3 SRR1039512 control   N052611 GSM1275866
## 4 SRR1039513 treated   N052611 GSM1275867
## 5 SRR1039516 control   N080611 GSM1275870
## 6 SRR1039517 treated   N080611 GSM1275871
```

Side note: Let's check the correspondance of the metadata nad count data setup. i.e. check if the first column of metadata is the same as the counts column headers

```
metadata[1]
```

```
##           id
## 1 SRR1039508
## 2 SRR1039509
## 3 SRR1039512
## 4 SRR1039513
## 5 SRR1039516
## 6 SRR1039517
## 7 SRR1039520
## 8 SRR1039521
```

```
colnames(counts)==metadata[1]
```

```
##           id
## [1,] TRUE
## [2,] TRUE
## [3,] TRUE
```

```
## [4,] TRUE
## [5,] TRUE
## [6,] TRUE
## [7,] TRUE
## [8,] TRUE
```

wrap the above code in all to tell us if all the outputs are true

```
all(colnames(counts)==metadata[1])
```

```
## [1] TRUE
```

##compare control to treated First we need to access all the control columns in our counts data. This is in the column “dex” using control.inds in [] gets that row of Trues and “1” gets the first col of this

```
control.inds <- metadata$dex=="control"
metadata[control.inds,1]
```

```
## [1] "SRR1039508" "SRR1039512" "SRR1039516" "SRR1039520"
```

Use these ids to access just the control columns of our ‘counts’ data

```
control.ids <- metadata[control.inds,]$id
control.ids
```

```
## [1] "SRR1039508" "SRR1039512" "SRR1039516" "SRR1039520"
```

```
control.mean <- rowMeans(counts[,control.ids])
head(control.mean)
```

```
## ENSG00000000003 ENSG00000000005 ENSG00000000419 ENSG00000000457 ENSG00000000460
##           900.75           0.00           520.50           339.75           97.25
## ENSG00000000938
##           0.75
```

##Do the same for the drug tested

```
treated.id <- metadata[metadata$dex == "treated",]$id
treated.mean <- rowMeans(counts[,treated.id])
```

We will combine our meancount data for bookkeeping purposes.

```
meancounts <- data.frame(control.mean, treated.mean)
```

Use the ‘ beside the 1 in the keyboard to also use code in this r,markdown
it will only show up in the r script: There are 38694 rows/genes in this dataset

##how many genes in this dataset


```
nrow(counts)
```

```
## [1] 38694
```

```
## Compare the control and treated
```

A quick plot of our progress so far

```
plot(meancounts)
```




This plot (above) needs to be altered to make it clearer. A log would be v useful

```
plot(meancounts, log = "xy")
```

```
## Warning in xy.coords(x, y, xlabel, ylabel, log): 15032 x values <= 0 omitted  
## from logarithmic plot
```

```
## Warning in xy.coords(x, y, xlabel, ylabel, log): 15281 y values <= 0 omitted  
## from logarithmic plot
```



We often use log transformations as they make life much easier...

```
log2(10/20)
```

```
## [1] -1
```

```
log2(80/20)
```

```
## [1] 2
```

```
meancounts$log2fc <- log2(meancounts[, "treated.mean"]/meancounts[, "control.mean"])
head(meancounts)
```

```
##           control.mean treated.mean    log2fc
## ENSG000000000003      900.75      658.00 -0.45303916
## ENSG000000000005         0.00         0.00         NaN
## ENSG0000000000419      520.50      546.00  0.06900279
## ENSG0000000000457      339.75      316.50 -0.10226805
## ENSG0000000000460       97.25       78.75 -0.30441833
## ENSG0000000000938         0.75         0.00      -Inf
```

Let's look for zeros in the meancounts

```
head(meancounts[,1:2])
```

```
##               control.mean treated.mean
## ENSG000000000003      900.75      658.00
## ENSG000000000005        0.00        0.00
## ENSG000000000419      520.50      546.00
## ENSG000000000457      339.75      316.50
## ENSG000000000460       97.25       78.75
## ENSG000000000938        0.75        0.00
```

```
head(meancounts[,1:2] == 0)
```

```
##               control.mean treated.mean
## ENSG000000000003      FALSE      FALSE
## ENSG000000000005       TRUE       TRUE
## ENSG000000000419      FALSE      FALSE
## ENSG000000000457      FALSE      FALSE
## ENSG000000000460      FALSE      FALSE
## ENSG000000000938      FALSE       TRUE
```

The which() function tells us the indices of TRUE entries in a logical vector

However, it is not useful in default mode on our type of multi column input...

```
inds <- which(meancounts[,1:2] == 0, arr.ind = TRUE)
head(inds)
```

```
##           row col
## ENSG000000000005    2  1
## ENSG000000004848   65  1
## ENSG000000004948   70  1
## ENSG000000005001   73  1
## ENSG000000006059  121  1
## ENSG000000006071  123  1
```

I only care about the rows here (if there is a zero in any column I will exclude this row eventually).

```
to.rm <- unique(sort(inds[, "row"]))
```

```
head(meancounts[to.rm,])
```

```
##               control.mean treated.mean log2fc
## ENSG000000000005        0.00        0.00   NaN
## ENSG000000000938        0.75        0.00 -Inf
## ENSG000000004848        0.00        0.25   Inf
## ENSG000000004948        0.00        0.00   NaN
## ENSG000000005001        0.00        0.00   NaN
## ENSG000000005102        1.00        0.00 -Inf
```

```
mycounts <- meancounts[-to.rm,]
head(mycounts)
```

```
##               control.mean treated.mean      log2fc
## ENSG000000000003      900.75      658.00 -0.45303916
## ENSG000000000419      520.50      546.00  0.06900279
## ENSG000000000457      339.75      316.50 -0.10226805
## ENSG000000000460       97.25       78.75 -0.30441833
## ENSG000000000971     5219.00     6687.50  0.35769358
## ENSG00000001036     2327.00     1785.75 -0.38194109
```

We now have 21817 genes remaining

```
nrow(mycounts)
```

```
## [1] 21817
```

How many of these genes are upregulated at the log2 fold-change threshold of +2 or greater

```
sum(mycounts$log2fc > 2)
```

```
## [1] 250
```

What is the percentage of this?

```
round((sum(mycounts$log2fc > 2) / nrow(mycounts)) *100,2)
```

```
## [1] 1.15
```

The calculate down regulated

```
sum(mycounts$log2fc < -2)
```

```
## [1] 367
```

```
round((sum(mycounts$log2fc < -2) / nrow(mycounts)) *100,2)
```

```
## [1] 1.68
```

#DESeq2 analysis We first need to set up the DESeq input object

```
dds <- DESeqDataSetFromMatrix(countData=counts,
                              colData=metadata,
                              design=~dex)
```

```
## converting counts to integer mode
```

```
## Warning in DESeqDataSet(se, design = design, ignoreRank): some variables in
## design formula are characters, converting to factors
```



```
dds
```

```
## class: DESeqDataSet
## dim: 38694 8
## metadata(1): version
## assays(1): counts
## rownames(38694): ENSG000000000003 ENSG000000000005 ... ENSG00000283120
## ENSG00000283123
## rowData names(0):
## colnames(8): SRR1039508 SRR1039509 ... SRR1039520 SRR1039521
## colData names(4): id dex celltype geo_id
```

Run the DESeq analysis pipeline

```
dds <- DESeq(dds)
```

```
## estimating size factors
```

```
## estimating dispersions
```

```
## gene-wise dispersion estimates
```

```
## mean-dispersion relationship
```

```
## final dispersion estimates
```

```
## fitting model and testing
```


```
res <- results(dds)
head(res)
```

```
## log2 fold change (MLE): dex treated vs control
## Wald test p-value: dex treated vs control
## DataFrame with 6 rows and 6 columns
##           baseMean log2FoldChange      lfcSE      stat      pvalue
##           <numeric>      <numeric> <numeric> <numeric> <numeric>
## ENSG000000000003 747.194195    -0.3507030  0.168246 -2.084470 0.0371175
## ENSG000000000005  0.000000         NA         NA         NA         NA
## ENSG000000000419 520.134160     0.2061078  0.101059  2.039475 0.0414026
## ENSG000000000457 322.664844     0.0245269  0.145145  0.168982 0.8658106
## ENSG000000000460  87.682625    -0.1471420  0.257007 -0.572521 0.5669691
## ENSG000000000938  0.319167    -1.7322890  3.493601 -0.495846 0.6200029
##           padj
##           <numeric>
## ENSG000000000003 0.163035
## ENSG000000000005      NA
## ENSG000000000419 0.176032
## ENSG000000000457 0.961694
## ENSG000000000460 0.815849
## ENSG000000000938      NA
```

A Volcano plot

This is a very common data viz of this type of data that does not really look like a volcano abline creates a line

```
plot(res$log2FoldChange, -log(res$padj))
abline(v=c(-2,2), col="gray")
abline(h=-log(0.05), col="gray")
```



##Adding annotation data We want to add meaningful gene names to our dataset so we can make some sense of what is going on here..

For this we will use two bioconductor packages, one does the work and is called **AnnotationDbi**. The other contains the data we are going to map between and is called **org.Hs.eg.db**

```
library("AnnotationDbi")
```

```
## Warning: package 'AnnotationDbi' was built under R version 4.1.2
```

```
library("org.Hs.eg.db")
```

```
##
```

```
columns(org.Hs.eg.db)
```

```
## [1] "ACCNUM"      "ALIAS"        "ENSEMBL"      "ENSEMBLPROT"  "ENSEMBLTRANS"
## [6] "ENTREZID"    "ENZYME"       "EVIDENCE"     "EVIDENCEALL"  "GENENAME"
## [11] "GENETYPE"    "GO"           "GOALL"        "IPI"           "MAP"
## [16] "OMIM"        "ONTOLOGY"     "ONTOLOGYALL"  "PATH"          "PFAM"
## [21] "PMID"        "PROSITE"      "REFSEQ"       "SYMBOL"        "UCSCCKG"
## [26] "UNIPROT"
```

Here we map to “SYMBOL” the common gene name that the world understands and wants

```
res$symbol <- mapIds(org.Hs.eg.db,
                     keys=row.names(res), # Our genenames
                     keytype="ENSEMBL",   # The format of our genenames
                     column="SYMBOL",     # The new format we want to add
                     multiVals="first")
```

```
## 'select()' returned 1:many mapping between keys and columns
```

```
head(res$symbol)
```

```
## ENSG000000000003 ENSG000000000005 ENSG000000000419 ENSG000000000457 ENSG000000000460
##      "TSPAN6"      "TNMD"      "DPM1"      "SCYL3"      "C1orf112"
## ENSG000000000938
##      "FGR"
```

```
head(res)
```

```
## log2 fold change (MLE): dex treated vs control
## Wald test p-value: dex treated vs control
## DataFrame with 6 rows and 7 columns
##      baseMean log2FoldChange    lfcSE    stat    pvalue
##      <numeric>    <numeric> <numeric> <numeric> <numeric>
## ENSG000000000003 747.194195   -0.3507030 0.168246 -2.084470 0.0371175
## ENSG000000000005  0.000000      NA      NA      NA      NA
## ENSG000000000419 520.134160    0.2061078 0.101059  2.039475 0.0414026
## ENSG000000000457 322.664844    0.0245269 0.145145  0.168982 0.8658106
## ENSG000000000460  87.682625   -0.1471420 0.257007 -0.572521 0.5669691
## ENSG000000000938  0.319167   -1.7322890 3.493601 -0.495846 0.6200029
##      padj      symbol
##      <numeric> <character>
## ENSG000000000003 0.163035      TSPAN6
## ENSG000000000005      NA      TNMD
## ENSG000000000419 0.176032      DPM1
## ENSG000000000457 0.961694      SCYL3
## ENSG000000000460 0.815849      C1orf112
## ENSG000000000938      NA      FGR
```

Lets save our results to date.

```
write.csv(res, file = "class15_allmyresults.csv")
```

##Pathway analysis

Let's try bring some insight back into this work. For this we will start with KEGG.

```
library(pathview)
library(gage)
library(gageData)

data(kegg.sets.hs)

# Examine the first 2 pathways in this kegg set for humans
head(kegg.sets.hs, 2)
```

```
## $'hsa00232 Caffeine metabolism'
## [1] "10"    "1544" "1548" "1549" "1553" "7498" "9"
##
## $'hsa00983 Drug metabolism - other enzymes'
## [1] "10"      "1066"  "10720" "10941" "151531" "1548"  "1549"  "1551"
## [9] "1553"    "1576"  "1577"  "1806"  "1807"  "1890"  "221223" "2990"
## [17] "3251"    "3614"  "3615"  "3704"  "51733" "54490" "54575" "54576"
## [25] "54577"    "54578" "54579" "54600" "54657" "54658" "54659" "54963"
## [33] "574537"  "64816" "7083"  "7084"  "7172"  "7363"  "7364"  "7365"
## [41] "7366"    "7367"  "7371"  "7372"  "7378"  "7498"  "79799" "83549"
## [49] "8824"    "8833"  "9"     "978"
```

Before we can use KEGG we need to get our gene identifiers in the correct format for KEGG, which is ENTREZ format in this case

```
head(rownames(res))
```

```
## [1] "ENSG000000000003" "ENSG000000000005" "ENSG000000000419" "ENSG000000000457"
## [5] "ENSG000000000460" "ENSG000000000938"
```

need to translate the rownames above into entrez format

```
head(columns(org.Hs.eg.db))
```

```
## [1] "ACCNUM"      "ALIAS"      "ENSEMBL"    "ENSEMBLPROT" "ENSEMBLTRANS"
## [6] "ENTREZID"
```

```
res$entrez <- mapIds(org.Hs.eg.db,
                     keys=row.names(res),
                     keytype="ENSEMBL",
                     column="ENTREZID",
                     multiVals="first")
```

```
## 'select()' returned 1:many mapping between keys and columns
```

The main `gage()` function requires a named vector of fold changes, where the names of the values are the Entrez gene IDs.

Note that we used the `mapIDs()` function above to obtain Entrez gene IDs (stored in `resentrez`) and we have the fold change results

```
foldchanges = res$log2FoldChange
head(foldchanges)
```

```
## [1] -0.35070302      NA  0.20610777  0.02452695 -0.14714205 -1.73228897
```

assign names to this vector that are the gene IDs that KEGG wants

```
names(foldchanges) <- res$entrez
head(foldchanges)
```

```
##          7105          64102          8813          57147          55732          2268
## -0.35070302      NA  0.20610777  0.02452695 -0.14714205 -1.73228897
```

Now we can pass this to the `gage()` function

```
#get the results
keggres = gage(foldchanges, gsets=kegg.sets.hs)
```

We can look at the `attributes()` of an object

```
attributes(keggres)
```

```
## $names
## [1] "greater" "less"  "stats"
```

```
head(keggres$less)
```

```
##                                     p.geomean stat.mean
## hsa05332 Graft-versus-host disease      0.0004250461 -3.473346
## hsa04940 Type I diabetes mellitus      0.0017820293 -3.002352
## hsa05310 Asthma                        0.0020045888 -3.009050
## hsa04672 Intestinal immune network for IgA production 0.0060434515 -2.560547
## hsa05330 Allograft rejection           0.0073678825 -2.501419
## hsa04340 Hedgehog signaling pathway    0.0133239547 -2.248547
##                                     p.val      q.val
## hsa05332 Graft-versus-host disease      0.0004250461 0.09053483
## hsa04940 Type I diabetes mellitus      0.0017820293 0.14232581
## hsa05310 Asthma                        0.0020045888 0.14232581
## hsa04672 Intestinal immune network for IgA production 0.0060434515 0.31387180
## hsa05330 Allograft rejection           0.0073678825 0.31387180
## hsa04340 Hedgehog signaling pathway    0.0133239547 0.47300039
##                                     set.size      exp1
## hsa05332 Graft-versus-host disease      40 0.0004250461
## hsa04940 Type I diabetes mellitus      42 0.0017820293
## hsa05310 Asthma                        29 0.0020045888
## hsa04672 Intestinal immune network for IgA production 47 0.0060434515
## hsa05330 Allograft rejection           36 0.0073678825
## hsa04340 Hedgehog signaling pathway    56 0.0133239547
```

The `pathview()` function will add our genes to a KEGG pathway as colored entries:

```
pathview(gene.data=foldchanges, pathway.id="hsa05310")
```

```
## 'select()' returned 1:1 mapping between keys and columns
```

```
## Info: Working in directory /Users/caitronabrennan/Documents/Bioinformatics_213/R Studio Class/bgg
```

```
## Info: Writing image file hsa05310.pathview.png
```

